

# Blog\_Churn

Snow

9/2/2021

## numerical\_df DF

### Define the question

I am a data science for a blogger. The question is making conclusion on who is likely to click on the ads in the blog and derive insights. Build a model that can predict if a person will click an ad or not based on the features in the dataframe ## Metric for success

In order to work on the above problem, you need to do the following:

- Define the question- the metric for success, the context, experimental design taken and the appropriateness of the available data to answer the given question.
- Find and deal with outliers, anomalies, and missing data within the dataset.
- Perform univariate and bivariate analysis.
- From your insights provide a conclusion and recommendation.
- Build a model using classification using decision trees and Support Vector Machine
- Get an accuracy => 80%

### Data Understanding (the context)

A Kenyan entrepreneur has created an online cryptography course and would want to advertise it on her blog. She currently targets audiences originating from various countries. In the past, she ran ads to advertise a related course on the same blog and collected data in the process. She would now like to employ your services as a Data Science Consultant to help her identify which individuals are most likely to click on her ads.

In order to work on the above problem, you need to do the following:

- Define the question, the metric for success, the context, experimental design taken and the appropriateness of the available data to answer the given question.
- Find and deal with outliers, anomalies, and missing data within the dataset.
- Perform univariate and bivariate analysis.
- From your insights provide a conclusion and recommendation.
- Build models and get the metrics

## Experimental design

1. Import the data to R
2. Perform data exploration
3. Define metrics for success
4. Perform Uni variate and Bivariate data Analysis
5. Provide conclusion

## Loading the current working directory

```
ad_df <- read.csv('advertising.csv', header = TRUE, sep = ',')
```

## Data Exploration

```
# Standardize column names with standard naming convention ie lowercase and replace spaces with '_'

# replace the spaces with underscores using gsub() function
names(ad_df) <- gsub(" ", "_", names(ad_df))

# lowercase
names(ad_df) <- tolower(names(ad_df))

# display the column names to confirm the changes
colnames(ad_df)

## [1] "daily.time.spent.on.site" "age"
## [3] "area.income"             "daily.internet.usage"
## [5] "ad.topic.line"           "city"
## [7] "male"                    "country"
## [9] "timestamp"               "clicked.on.ad"

# Preview dataset
head(ad_df)
```

```
##   daily.time.spent.on.site age area.income daily.internet.usage
## 1          68.95  35    61833.90          256.09
## 2          80.23  31    68441.85          193.77
## 3          69.47  26    59785.94          236.50
## 4          74.15  29    54806.18          245.89
## 5          68.37  35    73889.99          225.58
## 6          59.99  23    59761.56          226.74
##               ad.topic.line          city male  country
## 1   Cloned 5thgeneration orchestration Wrightburgh 0   Tunisia
## 2   Monitored national standardization   West Jodi 1     Nauru
## 3   Organic bottom-line service-desk     Davidton 0 San Marino
## 4 Triple-buffered reciprocal time-frame West Terrifurt 1     Italy
## 5       Robust logistical utilization   South Manuel 0    Iceland
## 6   Sharable client-driven software     Jamieberg 1     Norway
```

```
##           timestamp clicked.on.ad
## 1 2016-03-27 00:53:11           0
## 2 2016-04-04 01:39:02           0
## 3 2016-03-13 20:35:42           0
## 4 2016-01-10 02:31:19           0
## 5 2016-06-03 03:36:18           0
## 6 2016-05-19 14:30:17           0
```

```
# Finding the Shape of the dataset
dim(ad_df)
```

```
## [1] 1000  10
```

```
# Finding the datatypes of the data
str(ad_df)
```

```
## 'data.frame':  1000 obs. of  10 variables:
## $ daily.time.spent.on.site: num  69 80.2 69.5 74.2 68.4 ...
## $ age                     : int  35 31 26 29 35 23 33 48 30 20 ...
## $ area.income             : num  61834 68442 59786 54806 73890 ...
## $ daily.internet.usage    : num  256 194 236 246 226 ...
## $ ad.topic.line           : chr  "Cloned 5thgeneration orchestration" "Monitored national standardi
## $ city                    : chr  "Wrightburgh" "West Jodi" "Davidton" "West Terrifurt" ...
## $ male                    : int  0 1 0 1 0 1 0 1 1 1 ...
## $ country                 : chr  "Tunisia" "Nauru" "San Marino" "Italy" ...
## $ timestamp               : chr  "2016-03-27 00:53:11" "2016-04-04 01:39:02" "2016-03-13 20:35:42"
## $ clicked.on.ad           : int  0 0 0 0 0 0 0 1 0 0 ...
```

## Data cleaning

```
# checking for missing Data
colSums(is.na(ad_df))
```

```
## daily.time.spent.on.site      age      area.income
##                0                0                0
##   daily.internet.usage      ad.topic.line      city
##                0                0                0
##                male      country      timestamp
##                0                0                0
##      clicked.on.ad
##                0
```

There is no missing values in the dataset.

```
# Check for duplicated data in the ad_Df
ad_df1 <- ad_df[duplicated(ad_df),]
ad_df1
```

```
## [1] daily.time.spent.on.site age      area.income
```


```
## [4] daily.internet.usage      ad.topic.line      city
## [7] male                        country            timestamp
## [10] clicked.on.ad
## <0 rows> (or 0-length row.names)
```

There are no duplicated records in the dataset


```
str(ad_df)
```

```
## 'data.frame':    1000 obs. of  10 variables:
## $ daily.time.spent.on.site: num  69 80.2 69.5 74.2 68.4 ...
## $ age                     : int  35 31 26 29 35 23 33 48 30 20 ...
## $ area.income             : num  61834 68442 59786 54806 73890 ...
## $ daily.internet.usage    : num  256 194 236 246 226 ...
## $ ad.topic.line           : chr   "Cloned 5thgeneration orchestration" "Monitored national standardi
## $ city                    : chr   "Wrightburgh" "West Jodi" "Davidton" "West Terrifurt" ...
## $ male                    : int   0 1 0 1 0 1 0 1 1 1 ...
## $ country                 : chr   "Tunisia" "Nauru" "San Marino" "Italy" ...
## $ timestamp               : chr   "2016-03-27 00:53:11" "2016-04-04 01:39:02" "2016-03-13 20:35:42"
## $ clicked.on.ad           : int   0 0 0 0 0 0 0 1 0 0 ...
```

```
boxplot(ad_df$daily.time.spent.on.site, main = 'Daily Time Spent on-site')
```




```
boxplot(ad_df$age, main = 'age Boxplot')
```




```
boxplot(ad_df$area.income, main = 'Area Income Boxplot')
```

**Area Income Boxplot**



```
boxplot(ad_df$daily.internet.usage, main = 'Daily Internet usage boxplot')
```


## Daily Internet usage boxplot



From the boxplots, only the Area\_income column has outliers.

```
#Print out the outliers  
boxplot(ad_df$area.income, main = 'Area Income Boxplot')$out
```

## Area Income Boxplot



```
## [1] 17709.98 18819.34 15598.29 15879.10 14548.06 13996.50 14775.50 18368.57
```

There are outliers that do not look like they are in the extreme. There are areas where poverty is prevalent in such areas the total income could be that small.

```
str(ad_df)
```

```
## 'data.frame':    1000 obs. of  10 variables:
## $ daily.time.spent.on.site: num  69 80.2 69.5 74.2 68.4 ...
## $ age                     : int  35 31 26 29 35 23 33 48 30 20 ...
## $ area.income             : num  61834 68442 59786 54806 73890 ...
## $ daily.internet.usage    : num  256 194 236 246 226 ...
## $ ad.topic.line           : chr  "Cloned 5thgeneration orchestration" "Monitored national standardi
## $ city                    : chr  "Wrightburgh" "West Jodi" "Davidton" "West Terrifurt" ...
## $ male                    : int   0 1 0 1 0 1 0 1 1 1 ...
## $ country                 : chr  "Tunisia" "Nauru" "San Marino" "Italy" ...
## $ timestamp               : chr  "2016-03-27 00:53:11" "2016-04-04 01:39:02" "2016-03-13 20:35:42" ...
## $ clicked.on.ad           : int   0 0 0 0 0 0 0 1 0 0 ...
```

```
ad_df[['timestamp']] <- as.POSIXct(ad_df[['timestamp']],
                                   format = "%Y-%m-%d %H:%M:%S")
```

```
str(ad_df)
```

```
## 'data.frame':    1000 obs. of  10 variables:
```



```
## $ daily.time.spent.on.site: num 69 80.2 69.5 74.2 68.4 ...
## $ age : int 35 31 26 29 35 23 33 48 30 20 ...
## $ area.income : num 61834 68442 59786 54806 73890 ...
## $ daily.internet.usage : num 256 194 236 246 226 ...
## $ ad.topic.line : chr "Cloned 5thgeneration orchestration" "Monitored national standardi
## $ city : chr "Wrightburgh" "West Jodi" "Davidton" "West Terrifurt" ...
## $ male : int 0 1 0 1 0 1 0 1 1 1 ...
## $ country : chr "Tunisia" "Nauru" "San Marino" "Italy" ...
## $ timestamp : POSIXct, format: "2016-03-27 00:53:11" "2016-04-04 01:39:02" ...
## $ clicked.on.ad : int 0 0 0 0 0 0 0 1 0 0 ...
```

The timestamp column is now in the correct dtype

## Univariate Data Analysis

### Numerical Columns

```
summary(ad_df)
```

```
## daily.time.spent.on.site age area.income daily.internet.usage
## Min. :32.60 Min. :19.00 Min. :13996 Min. :104.8
## 1st Qu.:51.36 1st Qu.:29.00 1st Qu.:47032 1st Qu.:138.8
## Median :68.22 Median :35.00 Median :57012 Median :183.1
## Mean :65.00 Mean :36.01 Mean :55000 Mean :180.0
## 3rd Qu.:78.55 3rd Qu.:42.00 3rd Qu.:65471 3rd Qu.:218.8
## Max. :91.43 Max. :61.00 Max. :79485 Max. :270.0
## ad.topic.line city male country
## Length:1000 Length:1000 Min. :0.000 Length:1000
## Class :character Class :character 1st Qu.:0.000 Class :character
## Mode :character Mode :character Median :0.000 Mode :character
## Mean :0.481
## 3rd Qu.:1.000
## Max. :1.000
## timestamp clicked.on.ad
## Min. :2016-01-01 02:52:10 Min. :0.0
## 1st Qu.:2016-02-18 02:55:42 1st Qu.:0.0
## Median :2016-04-07 17:27:29 Median :0.5
## Mean :2016-04-10 10:56:04 Mean :0.5
## 3rd Qu.:2016-05-31 03:18:14 3rd Qu.:1.0
## Max. :2016-07-24 00:22:16 Max. :1.0
```

```
# Mean
mean.age <- mean(ad_df$age)
mean.age
```

```
age
```

```
## [1] 36.009
```

```
#median
median.age <- median (ad_df$age)
median.age
```

```
## [1] 35
```

```
# Function to get the mode.
getmode <- function(v) {
  uniqv <- unique(v)
  uniqv[which.max(tabulate(match(v, uniqv)))]
}
```

```
mode.age <- getmode(ad_df$age)
mode.age
```

```
## [1] 31
```

```
####Area income
```

```
mean.areaincome <- mean(ad_df$area.income)
mean.areaincome
```

```
## [1] 55000
```

```
median.areaincome <- median(ad_df$area.income)
median.areaincome
```


```
## [1] 57012.3
```

```
mode.areaincome <- getmode(ad_df$area.income)
mode.areaincome
```

```
## [1] 61833.9
```

```
hist(ad_df$area.income,
      main="Histogram for Area Income",
      xlab="Area income",
      border="blue",
      col="steelblue",)
```

## Histogram for Area Income



```
mean.daily.internet <- mean(ad_df$daily.internet.usage)
mean.daily.internet
```

**daily.internet.usage**

```
## [1] 180.0001
```


```
median.daily.internet <- median(ad_df$daily.internet.usage)
median.daily.internet
```

```
## [1] 183.13
```

```
mode.daily.internet <- getmode(ad_df$daily.internet.usage)
mode.daily.internet
```

```
## [1] 167.22
```

```
hist(ad_df$daily.internet.usage,
     main = 'Daily Intenet Usage',
     xlab="Daily internet Usage (mins)",
     border="blue",
     col="steelblue")
```



```
mean.dtsos <- mean(ad_df$daily.time.spent.on.site)
mean.dtsos
```

Daily time spent on site

```
## [1] 65.0002
```

```
median.dtsos <- median(ad_df$daily.time.spent.on.site)
median.dtsos
```

```
## [1] 68.215
```

```
mode.dtsos <- getmode(ad_df$daily.time.spent.on.site)
mode.dtsos
```

```
## [1] 62.26
```


```
uniq_clickers <- unique(ad_df$clicked.on.ad, )
length(uniq_clickers)
```

`clicked.on.ad`

```
## [1] 2
```

There are two categories of the people who clicked on ads Let us plot the frequency of each

```
clickers <- ad_df$clicked.on.ad
clickers_frequency <- table (clickers)
barplot(clickers_frequency, col = "steelblue")
```



There are 500 people who clicked on ads and another 500 did not click on the ads.

### Categorical Columns

```
####ad.topic.line
```

```
uniq_topic <- unique(ad_df$ad.topic.line, )
length(uniq_topic)
```

```
## [1] 1000
```

There are 1000 unique topic lines meaning it would be impossible to get a good visualization.

```
uniq_city <- unique(ad_df$city, )  
length(uniq_city)
```

city

```
## [1] 969
```

There are 969 unique cities hence it would also be impossible to get a good visualization

```
uniq_country <- unique(ad_df$country)  
length(uniq_country)
```

country

```
## [1] 237
```

There are 237 unique countries.

```
library(sf)
```

```
## Linking to GEOS 3.9.0, GDAL 3.2.1, PROJ 7.2.1
```

```
library(raster)
```

```
## Loading required package: sp
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:raster':
```

```
##
```

```
## intersect, select, union
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
library(spData)
```


```
## To access larger datasets in this package, install the spDataLarge  
## package with: 'install.packages('spDataLarge',  
## repos='https://nowosad.github.io/drat/', type='source')'
```

```
#library(spDataLarge)  
library(tmap)      # for static and interactive maps  
library(leaflet)   # for interactive maps  
library(ggplot2)
```

```
country <- ad_df$country  
countyfreq <- table(country)
```

## Gender

```
male <- ad_df$male  
male_freq <- table(male)  
barplot(male_freq, main= 'Gender Distribution', xlab="Gender",  
        ylab="Number of people",  
        border="red",  
        col="steelblue")
```



```
###Overall Summary
```

```
summary(ad_df)
```

```
##  daily.time.spent.on.site      age      area.income      daily.internet.usage
##  Min.   :32.60                Min.   :19.00      Min.   :13996      Min.   :104.8
##  1st Qu.:51.36                1st Qu.:29.00      1st Qu.:47032      1st Qu.:138.8
##  Median :68.22                Median :35.00      Median :57012      Median :183.1
##  Mean   :65.00                Mean   :36.01      Mean   :55000      Mean   :180.0
##  3rd Qu.:78.55                3rd Qu.:42.00      3rd Qu.:65471      3rd Qu.:218.8
##  Max.   :91.43                Max.   :61.00      Max.   :79485      Max.   :270.0
##  ad.topic.line      city      male      country
##  Length:1000      Length:1000      Min.   :0.000      Length:1000
##  Class :character      Class :character      1st Qu.:0.000      Class :character
##  Mode  :character      Mode  :character      Median :0.000      Mode  :character
##                                     Mean   :0.481
##                                     3rd Qu.:1.000
##                                     Max.   :1.000
##  timestamp      clicked.on.ad
##  Min.   :2016-01-01 02:52:10      Min.   :0.0
##  1st Qu.:2016-02-18 02:55:42      1st Qu.:0.0
##  Median :2016-04-07 17:27:29      Median :0.5
##  Mean   :2016-04-10 10:56:04      Mean   :0.5
##  3rd Qu.:2016-05-31 03:18:14      3rd Qu.:1.0
##  Max.   :2016-07-24 00:22:16      Max.   :1.0
```

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'

## The following objects are masked from 'package:raster':
##
## intersect, union

## The following objects are masked from 'package:base':
##
## date, intersect, setdiff, union
```

```
ad_df$Month_Yr <- format(as.Date(ad_df$timestamp), "%Y-%m")
head(ad_df)
```


```
##  daily.time.spent.on.site age area.income daily.internet.usage
##  1          68.95 35      61833.90          256.09
##  2          80.23 31      68441.85          193.77
##  3          69.47 26      59785.94          236.50
##  4          74.15 29      54806.18          245.89
##  5          68.37 35      73889.99          225.58
##  6          59.99 23      59761.56          226.74
##                                     ad.topic.line      city male      country
##  1      Cloned 5thgeneration orchestration      Wrightburgh      0      Tunisia
```



```
## 2    Monitored national standardization      West Jodi    1    Nauru
## 3      Organic bottom-line service-desk      Davidton    0 San Marino
## 4 Triple-buffered reciprocal time-frame West Terrifurt    1    Italy
## 5      Robust logistical utilization      South Manuel    0    Iceland
## 6      Sharable client-driven software      Jamieberg    1    Norway
##      timestamp clicked.on.ad Month_Yr
## 1 2016-03-27 00:53:11          0 2016-03
## 2 2016-04-04 01:39:02          0 2016-04
## 3 2016-03-13 20:35:42          0 2016-03
## 4 2016-01-10 02:31:19          0 2016-01
## 5 2016-06-03 03:36:18          0 2016-06
## 6 2016-05-19 14:30:17          0 2016-05
```

## Bivariate Analysis

```
ggplot(data = ad_df, mapping = aes(x = area.income)) +
  geom_freqpoly(mapping = aes(colour = clicked.on.ad), binwidth = 2000)
```



In areas where the income lies between 60,000 and & 70,000 there is a higher number of people clicking the ads ##### Correlation

```
#creating with only interger columns
numerical_df = ad_df[c("daily.time.spent.on.site", "age", "area.income", "daily.internet.usage", "male",
head(numerical_df)
```

```
##   daily.time.spent.on.site age area.income daily.internet.usage male
## 1                68.95  35    61833.90             256.09    0
## 2                80.23  31    68441.85             193.77    1
## 3                69.47  26    59785.94             236.50    0
## 4                74.15  29    54806.18             245.89    1
## 5                68.37  35    73889.99             225.58    0
## 6                59.99  23    59761.56             226.74    1
##   clicked.on.ad
## 1                0
## 2                0
## 3                0
## 4                0
## 5                0
## 6                0
```

```
correlation = cor(numerical_df)
correlation
```

```
##               daily.time.spent.on.site      age area.income
## daily.time.spent.on.site      1.00000000 -0.33151334  0.310954413
## age                          -0.33151334  1.00000000 -0.182604955
## area.income                   0.31095441 -0.18260496  1.000000000
## daily.internet.usage          0.51865848 -0.36720856  0.337495533
## male                         -0.01895085 -0.02104406  0.001322359
## clicked.on.ad                -0.74811656  0.49253127 -0.476254628
##               daily.internet.usage      male clicked.on.ad
## daily.time.spent.on.site      0.51865848 -0.018950855 -0.74811656
## age                          -0.36720856 -0.021044064  0.49253127
## area.income                   0.33749553  0.001322359 -0.47625463
## daily.internet.usage          1.00000000  0.028012326 -0.78653918
## male                         0.02801233  1.000000000 -0.03802747
## clicked.on.ad                -0.78653918 -0.038027466  1.00000000
```

```
library("PerformanceAnalytics")
```

```
## Loading required package: xts
```

```
## Loading required package: zoo
```

```
##
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':
##
##   as.Date, as.Date.numeric
```

```
##
## Attaching package: 'xts'
```

```
## The following object is masked from 'package:leaflet':
##
##   addLegend
```

```
## The following objects are masked from 'package:dplyr':
##
##   first, last


##
## Attaching package: 'PerformanceAnalytics'

## The following object is masked from 'package:graphics':
##
##   legend
```

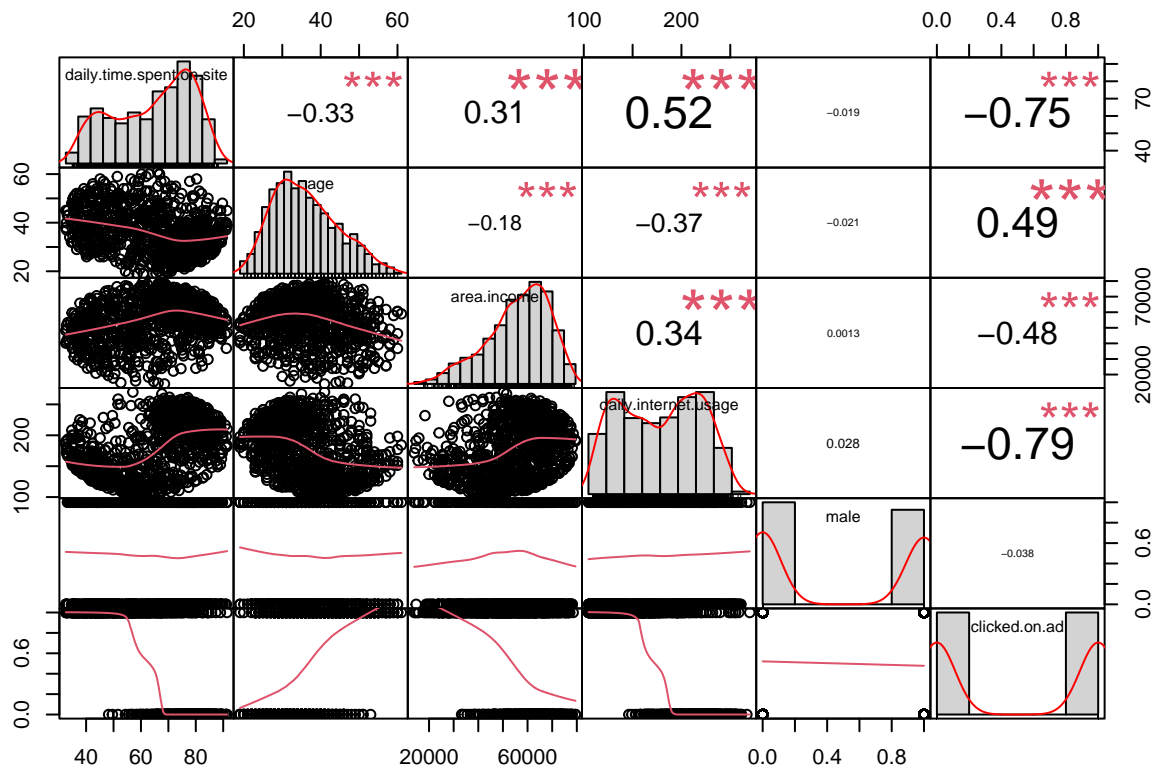
```
library(corrplot)
```

```
## corrplot 0.90 loaded
```

```
# Correlation Matrix
corrplot(correlation, method = 'number')
```



```
chart.Correlation(numerical_df, histogram = TRUE, pch = 19, )
```



The chart correlations gives a clear summary on the Bi-variate analysis of the dataframe.

## Bivariate Conclusion

From the analysis we can get several deductions: - daily.time.spent.on.site and the clicked.on.ad have an inverse. - The mean age of the population is 35, and as the age increased more people clicked on ads. - There is an inverse relationship between the daily time spent on site and the number of people who click the ads - There were slightly more females in the dataset. - The gender of the users had the least effect on the number of ads clicked and barely affected any other variables. '''

## Unsupervised learning.

```
# preview data structure
str(numerical_df)
```

```
## 'data.frame': 1000 obs. of 6 variables:
## $ daily.time.spent.on.site: num 69 80.2 69.5 74.2 68.4 ...
## $ age : int 35 31 26 29 35 23 33 48 30 20 ...
## $ area.income : num 61834 68442 59786 54806 73890 ...
## $ daily.internet.usage : num 256 194 236 246 226 ...
## $ male : int 0 1 0 1 0 1 0 1 1 1 ...
## $ clicked.on.ad : int 0 0 0 0 0 0 0 1 0 0 ...
```

```
head(numerical_df)
```

```
##   daily.time.spent.on.site age area.income daily.internet.usage male
## 1                68.95  35    61833.90           256.09    0
## 2                80.23  31    68441.85           193.77    1
## 3                69.47  26    59785.94           236.50    0
## 4                74.15  29    54806.18           245.89    1
## 5                68.37  35    73889.99           225.58    0
## 6                59.99  23    59761.56           226.74    1
##   clicked.on.ad
## 1             0
## 2             0
## 3             0
## 4             0
## 5             0
## 6             0
```

```
#We have to first make sure all the columns are in numerical format
numerical_df[,1:6] <- sapply(numerical_df[,1:6], as.numeric)
head(numerical_df)
```

```
##   daily.time.spent.on.site age area.income daily.internet.usage male
## 1                68.95  35    61833.90           256.09    0
## 2                80.23  31    68441.85           193.77    1
## 3                69.47  26    59785.94           236.50    0
## 4                74.15  29    54806.18           245.89    1
## 5                68.37  35    73889.99           225.58    0
## 6                59.99  23    59761.56           226.74    1
##   clicked.on.ad
## 1             0
## 2             0
## 3             0
## 4             0
## 5             0
## 6             0
```

```
# Normalizing the numerical variables of the data set. Normalizing the numerical values is really effective
# as it provides a measure from 0 to 1 which corresponds to min value to the max value of the data column
# We define a normal function which will normalize the set of values according to its minimum value and maximum value
normalize <- function(x) (
  return( ((x - min(x)) / (max(x) - min(x))) )
)
```

```
# Applying the normalization function
numerical_df$area.income<- normalize(numerical_df$area.income)
numerical_df$daily.internet.usage<- normalize(numerical_df$daily.internet.usage)
numerical_df$daily.time.spent.on.site<- normalize(numerical_df$daily.time.spent.on.site)
numerical_df$male<- normalize(numerical_df$male)
numerical_df$age<- normalize(numerical_df$age)
head(numerical_df)
```

```
##   daily.time.spent.on.site      age area.income daily.internet.usage male
```

```
## 1      0.6178820 0.3809524 0.7304725      0.9160310 0
## 2      0.8096209 0.2857143 0.8313752      0.5387456 1
## 3      0.6267211 0.1666667 0.6992003      0.7974331 0
## 4      0.7062723 0.2380952 0.6231599      0.8542802 1
## 5      0.6080231 0.3809524 0.9145678      0.7313234 0
## 6      0.4655788 0.0952381 0.6988280      0.7383460 1
## clicked.on.ad
## 1      0
## 2      0
## 3      0
## 4      0
## 5      0
## 6      0
```

This is a classification problem and for the models we will build two models the Decision trees model and the SVM model Lets begin

## Decision Tree

Importing the important libraries in modelling.

```
library(rpart,quietly = TRUE)
library(caret,quietly = TRUE)
library(rpart.plot,quietly = TRUE)
library(rpart.plot)
library(rattle)
```

```
## Loading required package: tibble
```

```
## Loading required package: bitops
```

```
## Rattle: A free graphical interface for data science with R.
## Version 5.4.0 Copyright (c) 2006-2020 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(e1071)
```

```
##
```

```
## Attaching package: 'e1071'
```

```
## The following objects are masked from 'package:PerformanceAnalytics':
```

```
##
```

```
##      kurtosis, skewness
```

```
## The following object is masked from 'package:raster':
```

```
##
```

```
##      interpolate
```

Start with data preparation such as Splitting the data into training and testing set

```

set.seed(42)
dttrain <- sample(1:nrow(numerical_df),size = ceiling(0.80*nrow(numerical_df)),replace = FALSE)
# training set
dt_train <- numerical_df[dttrain,]
# test set
dt_test <- numerical_df[-dttrain,]

```

we performed an 80-20 split on the data

```

# we are defining the penalty matrix for the decision tree to ensure that the model has more accurate p
# The penalty will multiply an error by 10
# Penalty matrix
penalty.matrix <- matrix(c(0, 1, 10,0), byrow = TRUE, nrow = 2)

```

```

dtree <- rpart(clicked.on.ad ~., data = dt_train, parms=list(loss=penalty.matrix), method = 'class')
dtree

```

```

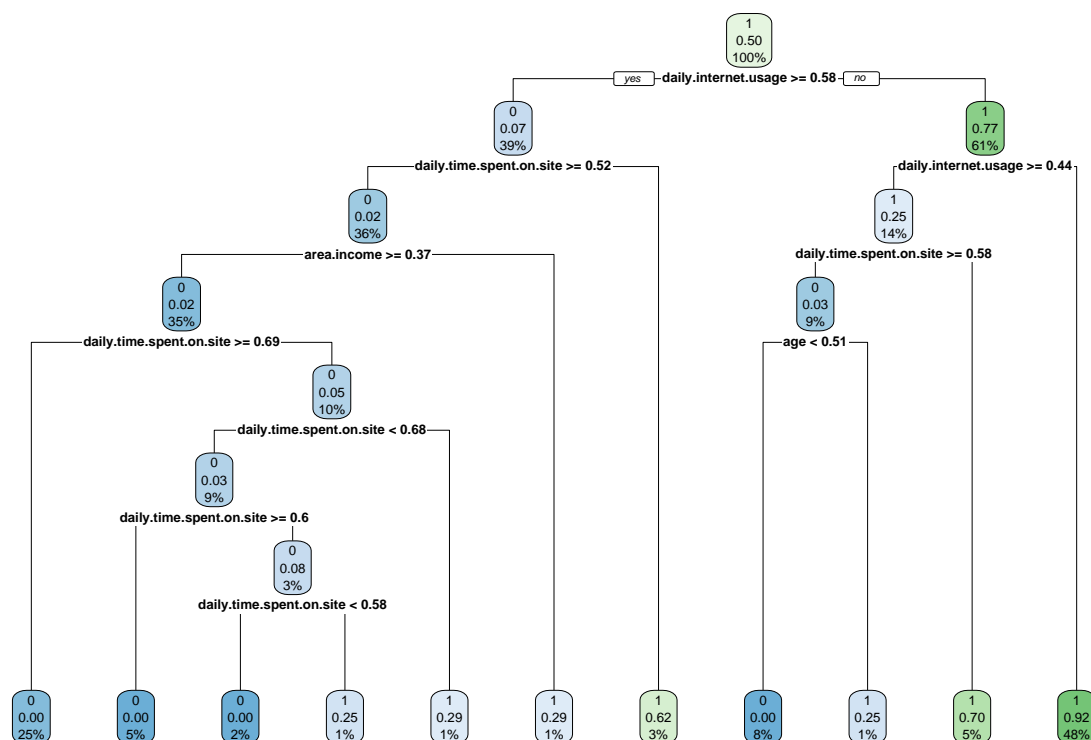
## n= 800
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
##  1) root 800 399 1 (0.498750000 0.501250000)
##    2) daily.internet.usage>=0.5780058 309 220 0 (0.928802589 0.071197411)
##      4) daily.time.spent.on.site>=0.5224375 285 70 0 (0.975438596 0.024561404)
##        8) area.income>=0.3704821 278 50 0 (0.982014388 0.017985612)
##          16) daily.time.spent.on.site>=0.690804 202 10 0 (0.995049505 0.004950495) *
##            17) daily.time.spent.on.site< 0.690804 76 40 0 (0.947368421 0.052631579)
##              34) daily.time.spent.on.site< 0.6808601 69 20 0 (0.971014493 0.028985507)
##                68) daily.time.spent.on.site>=0.6013089 43 0 0 (1.000000000 0.000000000) *
##                  69) daily.time.spent.on.site< 0.6013089 26 20 0 (0.923076923 0.076923077)
##                    138) daily.time.spent.on.site< 0.5802312 18 0 0 (1.000000000 0.000000000) *
##                      139) daily.time.spent.on.site>=0.5802312 8 6 1 (0.750000000 0.250000000) *
##                        35) daily.time.spent.on.site>=0.6808601 7 5 1 (0.714285714 0.285714286) *
##                          9) area.income< 0.3704821 7 5 1 (0.714285714 0.285714286) *
##                            5) daily.time.spent.on.site< 0.5224375 24 9 1 (0.375000000 0.625000000) *
##                              3) daily.internet.usage< 0.5780058 491 112 1 (0.228105906 0.771894094)
##                                6) daily.internet.usage>=0.4388243 111 83 1 (0.747747748 0.252252252)
##                                  12) daily.time.spent.on.site>=0.581506 74 20 0 (0.972972973 0.027027027)
##                                    24) age< 0.5119048 66 0 0 (1.000000000 0.000000000) *
##                                      25) age>=0.5119048 8 6 1 (0.750000000 0.250000000) *
##                                        13) daily.time.spent.on.site< 0.581506 37 11 1 (0.297297297 0.702702703) *
##                                          7) daily.internet.usage< 0.4388243 380 29 1 (0.076315789 0.923684211) *

```

```

rpart.plot(dtree)

```



```
# Calculating the metrics of the decision Tree model
# Predictions Dtree model
predt <- predict(object = dtree, dt_test[,-6], type = 'class')
#calculating accuracy
t <- table(dt_test$clicked.on.ad, predt)
confusionMatrix(t)
```

```
## Confusion Matrix and Statistics
##
##      predt
##      0  1
## 0  81  20
## 1   2  97
##
##               Accuracy : 0.89
##               95% CI : (0.8382, 0.9298)
##      No Information Rate : 0.585
##      P-Value [Acc > NIR] : < 2.2e-16
##
##               Kappa : 0.7804
##
##  Mcnemar's Test P-Value : 0.0002896
##
##               Sensitivity : 0.9759
##               Specificity : 0.8291
##               Pos Pred Value : 0.8020
```



```
##          Neg Pred Value : 0.9798
##          Prevalence : 0.4150
##          Detection Rate : 0.4050
##          Detection Prevalence : 0.5050
##          Balanced Accuracy : 0.9025
##
##          'Positive' Class : 0
##
```

The decision tree model has an accuracy of 89% That is quite acceptable as the metrics for success needed an accuracy of 80%

## LinearSVM

```
library('caret')
# Performing an 80 - 20 split

svmtrain <- createDataPartition(y = numerical_df$clicked.on.ad, p= 0.8, list = FALSE)
training <- numerical_df[svmtrain,]

testing <- numerical_df[-svmtrain,]
```

```
# Preview the dimensions of the training and testing data
dim(training)
```

```
## [1] 800  6
```

```
dim(testing)
```

```
## [1] 200  6
```

```
# Building a LinearSVM model
# SVM model
```

```
classifierL = svm(formula = clicked.on.ad ~ .,
                  data = training,
                  type = 'C-classification',
                  kernel = 'linear')
```

```
# Running the metrics for the linear classification svm
y_predL = predict(classifierL, newdata = testing)
```

```
cmL = table(testing$clicked.on.ad, y_predL)
confusionMatrix(cmL)
```

```
## Confusion Matrix and Statistics
```

```
##
##      y_predL
##      0  1
## 0 98  2
```

```
## 1 2 98
##
##          Accuracy : 0.98
##          95% CI : (0.9496, 0.9945)
##    No Information Rate : 0.5
##    P-Value [Acc > NIR] : <2e-16
##
##          Kappa : 0.96
##
## Mcnemar's Test P-Value : 1
##
##          Sensitivity : 0.98
##          Specificity : 0.98
##    Pos Pred Value : 0.98
##    Neg Pred Value : 0.98
##          Prevalence : 0.50
##    Detection Rate : 0.49
##    Detection Prevalence : 0.50
##    Balanced Accuracy : 0.98
##
##    'Positive' Class : 0
##
```

The Linear SVM model has an accuracy of 98% which is quite an improvement from the Decision tree model. However, we can still challenge the model, to find a better model that might account for overfitting.

## Challenging the solution

```
# Running the classifier with a
classifierRB = svm(formula = clicked.on.ad ~ .,
                   data = training,
                   type = 'C-classification',
                   kernel = 'sigmoid')

# Running the metrics for the linear classification sum
y_predRB = predict(classifierRB, newdata = testing)

cmRB = table(testing$clicked.on.ad, y_predRB)
confusionMatrix(cmRB)
```

```
## Confusion Matrix and Statistics
##
##      y_predRB
##      0  1
## 0 93  7
## 1  4 96
##
##          Accuracy : 0.945
##          95% CI : (0.9037, 0.9722)
##    No Information Rate : 0.515
##    P-Value [Acc > NIR] : <2e-16
```

```
##
##           Kappa : 0.89
##
## Mcnemar's Test P-Value : 0.5465
##
##           Sensitivity : 0.9588
##           Specificity : 0.9320
##           Pos Pred Value : 0.9300
##           Neg Pred Value : 0.9600
##           Prevalence : 0.4850
##           Detection Rate : 0.4650
##           Detection Prevalence : 0.5000
##           Balanced Accuracy : 0.9454
##
##           'Positive' Class : 0
##
```

With a sigmoid kernel on SVM technique the accuracy is (94.5%). This might be a better model because it would account for the overfitting.

## Conclusion

Was the data enough to answer the given questions? The data provided was sufficient in the analysis.

Do you have any recommendation on what data should be added? I think that the data was good, however, having more data would not be bad and training a bigger dataframe could lead to more accurate models.

The data was cleaned and used for analysis in the data frame